

Project Report: eBay Policy Assistant – RAG Chatbot with Streaming Responses

1. Document Structure & Chunking Logic

The provided input document, `AI_Training_Document.pdf`, contains over 10,000 words of policy-related content, including eBay's Terms and Conditions and User Agreements. The document was first cleaned to remove repetitive headers, footers, and irrelevant formatting using PyMuPDF (`fitz`). Specifically, lines starting with phrases like "eBay", "Page", or empty lines were filtered out. The cleaned output was saved as `cleaned_text.txt`.

To ensure meaningful context preservation for the chatbot, the text was segmented using **sentence-aware chunking** via `spaCy`. Each chunk was capped at approximately **150 words**, maintaining sentence boundaries to preserve semantic integrity. The result, `document_chunks.txt`, was used for embedding and retrieval.

2. Embedding Model & Vector Database

To semantically index the policy content, the model **BAAI/bge-small-en** from Hugging Face was employed. This model is lightweight yet highly performant for sentence embeddings and supports robust semantic similarity.

The embeddings were computed for all chunks and stored using the **FAISS** library, a fast similarity search library. FAISS enables efficient **vector-based retrieval** of the most relevant document segments based on a user's question. The resulting files:

- `faiss_index.idx`: Vector index file.
- `chunks.pkl`: Serialized chunk metadata.

These files are stored in the `/vectordb/` directory and are loaded at runtime by the RAG pipeline.

3. Prompt Format & Response Generation Logic

The Retrieval-Augmented Generation (RAG) pipeline involves two stages:

a. Retriever:

Top $k=3$ most relevant chunks are retrieved from the FAISS index based on cosine similarity with the query embedding.

b. Prompt Construction:

Retrieved chunks are injected into a conversational prompt template as "context" with the following instruction to the LLM:

“You are a smart and helpful AI assistant... Use the context below to answer the user's question, but do not mention the context explicitly... If the context does not contain enough information, simply say: 'I don't know.'”

The prompt is passed to the **Meta-LLaMA-3-8B-Instruct** model hosted via the **Hugging Face Inference API**, with API key access securely managed using a `.env` file.

c. Generator:

The response is generated in **real-time** and **streamed token-by-token** in the chatbot interface. This simulates an interactive assistant and improves perceived response speed.

4. Example Queries and Responses

This section presents example interactions with the RAG chatbot, demonstrating both successful and failed cases. Each example includes the user query, the response from the chatbot, and the analysis of retrieved context.

Failure Cases and Improvements

In the initial phase of development, we encountered some failure responses when testing certain queries. These failures were due to suboptimal prompt structure, higher chunk size, and fallback-based logic.

Failure Cases Using Older Version of RAGPipeline

Failure Code Characteristics:

- Chunk size used: ~**300 words**
- Prompt design was restrictive
- Prompt lacked flexibility and guidance for natural language generation

Examples of Failure Cases

User Question: What happens if I list items in the wrong category on eBay?

Bot Response: *"Sorry, I don't have enough information to respond clearly. Maybe ask it another way?"*

User Question: Can I use eBay to sell illegal items?

Bot Response: *"I couldn't find a specific answer right now. Try asking about something else!"*

Improvements Made

- **Reduced Chunk Size:**
 - *Before:* Each chunk contained ~300 words.
 - *After:* Reduced to ~150 words for more focused and precise retrieval.
- **Improved Prompt Logic:**
 - *Before:* Required query keywords to be explicitly present in retrieved chunks to proceed.
 - *After:* Allowed inference based on semantic relevance, enabling better and more flexible responses.
- **Enhanced Prompt Style:**
 - *Before:* Prompt was rigid and often led to fallback if the answer wasn't exact.
 - *After:* Prompt became more conversational and user-friendly, encouraging the use of bullet points and numbered lists for clarity.
- **Better Handling of Model Responses:**
 - *Before:* The model defaulted to generic fallback responses when unsure.
 - *After:* The model was guided to generate confident, helpful answers even with partial information, reducing fallback triggers.

Why the Change in Chunk Size?

Reducing the chunk size improved retrieval precision by narrowing the semantic scope of each chunk. This made the retrieved passages more focused and contextually aligned with user queries, which significantly increased response relevance.

Success Cases After Improvements

User Question: What happens if I list items in the wrong category on eBay?

Bot Response:

If you list an item in the wrong category on eBay, it may affect its visibility and searchability. Additionally, it's possible that your listing may not be immediately searchable by keyword or category, which could impact its discoverability. It's also important to note that you, as the seller, are responsible for ensuring the accuracy and content of your listing, including the category it's listed in. If you're unsure about the correct category for your item, you can consult eBay's category guidelines or reach out to eBay's customer support for assistance.

Sources:

1. If you wish to dispute the information a collection agency reported to a credit bureau regarding your eBay account...
2. Sellers and buyers are responsible for complying with all laws and regulations...
3. Such information may be based on the aggregated sales and performance history of similar sold and/or current listings...

User Question: Can I use eBay to sell illegal items?

Bot Response:

No, you cannot use eBay to sell illegal items. eBay's User Agreement explicitly prohibits breaching or circumventing any laws, regulations, or third-party rights. Additionally, eBay has a policy against selling items that are illegal or violate any applicable laws or regulations.

Sources:

1. Using eBay In connection with using or accessing our Services you agree to comply with this User Agreement...
2. If we believe you are violating this User Agreement or any of our policies...
3. If we believe you are violating our policies prohibiting Offers to buy or sell outside of eBay...

User Question: What is eBay?

Bot Response:

eBay is a platform that offers access to and use of various services, including mobile applications, tools, and online marketplaces. It's a digital platform where individuals and businesses can buy and sell a wide range of products and services, including vehicles.

Sources:

1. User Agreement 1. Introduction This User Agreement, the Mobile Application Terms of Use, and all policies and additional terms posted on and in our sites, applications, tools, and services (collectively "Services") set out the terms on which eBay offers you access to and use of our Services. You ...
2. WITHOUT LIMITING THE GENERALITY OF THE FOREGOING, WE (INCLUDING OUR PARENT, SUBSIDIARIES, AND AFFILIATES, AND OUR AND THEIR OFFICERS,

DIRECTORS, AGENTS, AND EMPLOYEES) ARE NOT LIABLE, AND YOU AGREE NOT TO HOLD US RESPONSIBLE, FOR ANY DAMAGES OR LOSSES (INCLUDING, BUT NOT LIMITED TO, LOSS OF MONEY...

3. eBay does not maintain an inventory of vehicles for sale and does not sell, exchange, buy, offer for sale, negotiate or attempt to negotiate sales or exchanges of an interest in any vehicle listed on eBay sites. eBay has never: (1) held title for any vehicle listed on eBay sites; (2) inspected an...

User Question: Does eBay provide insurance or guarantee for vehicle purchases?

Bot Response:

No, eBay does not provide insurance or guarantee for vehicle purchases. According to the text, neither the financing of nor insurance relevant to vehicles is controlled or guaranteed by eBay. You are advised to seek professional expert assistance and advice in these areas, as eBay is not an expert on matters such as vehicle insurance or guarantee.

Sources:

1. IN ADDITION TO, AND WITHOUT LIMITING THE GENERALITY OF THE DISCLAIMERS OF WARRANTIES DETAILED IN SECTION 16, YOU AGREE THAT: (A) EBAY DOES NOT WARRANT AGAINST DEFECTS...
2. Any discrepancies or mistakes made regarding vehicle availability, quality, safety, condition, pricing, financing, insurance and the like are not the fault...
3. eBay does not maintain an inventory of vehicles for sale and does not sell, exchange, buy, offer for sale, negotiate or attempt to negotiate sales...

User Question: How does eBay handle duplicate listings?

Bot Response:

eBay has policies in place to ensure a seamless buying and selling experience. When it comes to duplicate listings, eBay takes steps to prevent and address them. Here's how it works:

- If you accidentally list the same item multiple times, eBay may merge the listings or remove duplicates to prevent confusion and ensure that buyers can find what they're looking for.
- In some cases, eBay may also flag or remove duplicate listings that are intentionally posted to manipulate prices or create confusion.
- To avoid duplicate listings, it's essential to review your listings before posting and ensure you're not duplicating an existing listing.

Sources:

1. Such information may be based on the aggregated sales and performance history of similar sold and/or current listings...
2. eBay can't guarantee exact listing duration; your fixed-price listings may renew automatically...
3. Some actions that may constitute infringement are reproducing, copying, or distributing content that belongs to eBay or others...

User Question: Can eBay guarantee that a seller will complete a transaction?

Bot Response:

According to the provided information, eBay does not guarantee that a seller will complete a transaction. In fact, it's explicitly stated that "eBay has no control over and does not guarantee: the ability of sellers to sell items..." This means that eBay cannot ensure that a seller will indeed complete a transaction or fulfill their obligations.

Sources:

1. Most eBay sales go smoothly, but if there's a problem... buyers can get their money back...
2. Payments received by the eBay Payment Entity from you will satisfy your obligations...
3. We may use AI-powered tools to support fraud detection... availability and accuracy of these tools are not guaranteed...

User Question: How does eBay handle unauthorized use of an account? Step by step.

Bot Response:

eBay has a strict policy against unauthorized use of accounts. If they suspect or detect unauthorized access, they will take immediate action. Here's a step-by-step overview:

1. **Detection:** Suspicious activity is detected via login history, IP logs, etc.
2. **Investigation:** eBay confirms unauthorized access.
3. **Account Suspension:** Account access is temporarily locked.
4. **Notification:** User is informed of the issue.
5. **Verification:** User must verify identity.
6. **Account Termination:** If verification fails, the account may be closed.
7. **Legal Action:** In serious cases, law enforcement may be involved.
8. **Prevention:** Guidance is provided to improve account security.

Sources:

1. If you require that authorized third parties... we offer a Multi-User Account Access program...
2. If we believe you're violating policies... we may limit or suspend your account...
3. Using eBay requires you to comply with all applicable laws and policies...

5. Notes on Hallucinations, Limitations, and Performance

Strengths:

- Real-time streaming provides a highly responsive user experience.
- Compact architecture with clear modularity (retriever, generator, UI).
- **Fast Response Time:** The Meta-LLaMA-3-8B-Instruct model, when accessed via the Hugging Face Inference API (featherless-ai provider), delivers responses with very low latency, enabling real-time token-level streaming in the Streamlit chatbot interface.

- **Improved Response Quality via Prompt Engineering:** Clear and conversational prompt formatting using a role-based system prompt significantly improved the quality of model answers, making them more coherent and user-friendly.
- **More Accurate Retrieval with Smaller Chunk Size:** Reducing chunk size from 300 words to ~150 words enhanced retrieval precision. The smaller, semantically focused chunks ensured that only the most relevant context is passed to the LLM.
- **Well-Structured Multi-Step Query Handling:** The model was able to handle complex, multi-step questions effectively when guided with structured prompts and dense context.
- **No External Hallucination:** The final version prevents hallucination by instructing the model to strictly rely on the given context, and to respond with "I don't know" when the answer is not present—ensuring factual accuracy.

Limitations:

- **Hallucinations:** Rare but possible if the retrieved chunks are sparse or irrelevant.
- **No fine-tuning:** The model relies solely on prompt engineering.
- **API dependency:** Uses Hugging Face Inference API with Meta LLaMA 3, which may introduce latency or availability limits depending on usage tier.

Conclusion

This RAG chatbot provides a streamlined solution to query lengthy policy documents using retrieval and generation techniques. It is ideal for FAQ-style automation or internal policy assistants, and can be extended with memory, multi-document support, or role-based responses.

Project Resources & Demonstration

- **GitHub Repository:**
<https://github.com/TinaSingh19/rag-chatbot-ebay-policy>
- **Demo Video Folder:**
[Google Drive - Demo Videos \(Sample Queries + Streaming Responses\)](#)